ICSIC 2026 — Submission 2270

Changes made after review

AAURA: An AI-Assisted Campus Engagement Platform

This table lists the changes made in the revised manuscript.

What I changed	The change I made
The prediction claim	The paper no longer treats the synthetic result as evidence of real AAUP attendance. It states that no labeled campus outcomes exist yet, that the numbers measure recovery of a known rule, and that the score only advises staff. It does not publish or cancel an event. The discussion names the next study: retrain on one semester of stored check-ins and test the same metrics on the next semester.
The synthetic data generator	Section V-A now gives the feature ranges, the success formula, the bonuses and penalties, the 0.5 label cutoff, and the 6% label flip. It reports that flipped labels put a ceiling near 0.61 on low-success recall. Regenerating the data is also reported: macro F1 is 0.88 with no label noise, 0.67 at 12% noise, and 0.61 at 20% noise. Halving the match weights raises macro F1 to 0.83; multiplying them by 1.5 lowers it to 0.64.
What the model can see	A new subsection separates draft mode from live mode. Before publication the model uses only what the organizer typed. After publication it uses reservations and profiles that already exist. Feedback is kept for a later label and is not a model input. A promotion factor applied by the API is left out of the reported numbers.
The train and test split	The hyperparameters were fixed before the test set was opened. The 20% test split (seed 42) is scored once. Cross-validation, calibration, and the decision threshold use only the training split. The earlier cross-validation on the full table was replaced by training-split figures: accuracy 0.905 ± 0.006 and macro F1 0.749 ± 0.019.
Baselines	Accuracy is no longer the headline. Table II adds a majority-class baseline (accuracy 0.863, macro F1 0.46), a hand-written low-success rule, and logistic regression with and without class weights. The noise-free generator rule is reported as a ceiling, not as a competing system.
Low-success events, thresholds, and calibration	Low-success is the class used for recall and precision-recall area. Table II reports macro F1, low-success recall, and PR-AUC for the deployed model, class-weighted boosting, a threshold of 0.58 chosen on the training folds, and a calibrated model. Class weights raise low-success recall from 0.44 to 0.56 and lower precision from 0.66 to 0.51. Calibration lowers the probability error from 0.031 to 0.013. Fig. 2 adds the precision-recall curve and the reliability diagram. Fig. 3 adds the confusion matrices and the threshold sweep.
Match features and the display bands	On the test split, High, Medium, and Low separate observed success (0.93, 0.71, and 0.34). The display score ranks events slightly worse than the probability alone (ROC area 0.76 against 0.78), so it is kept only as a presentation aid. Removing interest and skill match drops macro F1 from 0.73 to 0.56 and the precision-recall area from 0.55 to 0.20.
Shams evaluation	Section IV-B adds 120 annotated messages (45 English, 45 Arabic, 22 mixed, and 8 Arabizi), labeled by two annotators. Table I gives precision, recall, and F1 for each field. Against a baseline that accepts only the English name of a label, profile F1 rises from 0.37 to 0.60. Removing the Arabic aliases drops Arabic profile F1 from 0.43 to 0.16. Confidence tracks how many items were extracted (Spearman 0.76) much more than whether the profile is correct (Spearman 0.27), so it is not described as the probability of a correct profile.
Matching events to students	Section VI writes the on-device score: 0.42 times interest overlap, 0.38 times skill overlap, and 0.12 times overlap with the organizer's target interests, plus a major bonus. A suggestion is shown only at 0.12 or above, and at most eight events are shown. Topics outside the profile are dropped. Ranking the same events with the gold profile and with the Shams profile gives suggestion lists that overlap with Jaccard 0.59. Across 4,160 student-event pairs, profiling errors change the success probability by 0.038 on average and change the band in 3.1% of pairs.
Wording	The sentences "It is not an LLM" and "We do not make up a score" were rewritten. Section II-C now says Shams is a lexicon and pattern matcher, not a generator, and that questions outside profiling are refused. The system section says the API returns a failure when the model cannot be reached and does not fill in a score. The introduction states three contributions tied to the new evaluation.

ICSIC 2026 — Submission 2270

Changes made after review

AAURA: An AI-Assisted Campus Engagement Platform

What I changed	The change I made
Figures	The evaluation plots requested in the reviews were added as Fig. 2 and Fig. 3. The application screens were reduced to one figure, Fig. 1.